

Benjamin Ye
CS/CNS/EE 156a: Learning Systems (Fall 2023)
November 20, 2023

Homework 8

Problem	Answer
1	[d]
2	[a]
3	[a]
4	[c]
5	[d]
6	[b]
7	[b]
8	[c]
9	[e]
10	[c]

Primal Versus Dual Problem

1. Recall that N is the size of the data set and d is the dimensionality of the input space. The original formulation of the hard-margin SVM problem (minimize $\mathbf{w}^T \mathbf{w} / 2$ subject to the inequality constants), without going through the Lagrangian dual problem, is

Answer: [d] a quadratic programming problem with $d + 1$ variables

In the hard-margin SVM problem, we minimize $\mathbf{w}^T \mathbf{w} / 2$ using the constraint $y_n(\mathbf{w}^T x_n + b) \geq 1$. Since the only unknowns in the constraint are \mathbf{w} and b , which have dimensions of d and 1, respectively, the hard-margin SVM problem is a quadratic programming problem with $d + 1$ variables.

SVM with Soft Margins

Notice: The following problems deal with a real-life data set. In addition, the computational packages you use may employ different heuristics and require different tweaks. This is a typical situation that a machine learning (ML) practitioner faces. There are uncertainties, and the answers may or may not match our expectations. Although this situation is not as “sanitized” as other homework problems, it is important to go through it as part of the learning experience.

In the rest of the problems of this homework set, we apply soft-margin SVM to the handwritten digits from the processed U.S. Postal Service Zip Code data set. Download the data (extracted features of intensity and symmetry) for training and testing:

<http://www.amlbook.com/data/zip/features.train>

<http://www.amlbook.com/data/zip/features.test>

(The format of each row is: **digit**, **intensity**, **symmetry**.) We will train two types of binary classifiers; one-versus-one (one digit is class +1 and another digit is class −1, with the rest of the digits disregarded), and one-versus-all (one digit is class +1 and the rest of the digits are class −1).

The data set has thousands of points, and some quadratic programming packages cannot handle this size. You may need stronger SVM packages such as **fitcsvm** in MATLAB or the free **libsvm**.

Implement SVM with soft margins on the above zip code data set by solving

$$\min_{\alpha} \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N \alpha_n \alpha_m y_n y_m K(\mathbf{x}_n, \mathbf{x}_m) - \sum_{n=1}^N \alpha_n \quad \text{s.t.} \quad \sum_{n=1}^N y_n \alpha_n = 0, \quad 0 \leq \alpha_n \leq C \quad \text{for } n = 1, \dots, N$$

When evaluating E_{in} and E_{out} of the resulting classifier, use binary classification error. E_{out} is estimated using the test set.

Practical remarks:

- (i) For this homework, do not scale the data when you use **libsvm** or other packages or you may inadvertently change the (effective) kernel and get different results.
- (ii) In some packages, you need to specify double precision.
- (iii) In 10-fold cross validation, if the data size is not a multiple of 10, the sizes of the 10 subsets may be off by 1 data point.
- (iv) Some packages have software parameters whose values affect the outcome. ML practitioners must deal with this kind of added uncertainty.

Polynomial Kernels

Consider the polynomial kernel $K(\mathbf{x}_n, \mathbf{x}_m) = (1 + \mathbf{x}_n^T \mathbf{x}_m)^Q$, where Q is the degree of polynomial.

2. With $C = 0.01$ and $Q = 2$, which of the following classifiers has the highest E_{in} ?

Answer: [a] 0 versus all

3. With $C = 0.01$ and $Q = 2$, which of the following classifiers has the lowest E_{in} ?

Answer: [a] 1 versus all

4. Comparing the two selected classifiers from Problems 2 and 3, which of the following values is the closest to the difference between the number of support vectors of these two classifiers?

Answer: [c] 1,800

5. Consider the 1 versus 5 classifier with $Q = 2$ and $C \in \{0.001, 0.01, 0.1, 1\}$. Which of the following statements is correct? Going up or down means strictly so.

Answer: [d] Maximum C achieves the lowest E_{in} .

6. In the 1 versus 5 classifier, comparing $Q = 2$ with $Q = 5$, which of the following statements is correct?

Answer: [b] When $C = 0.001$, the number of support vectors is lower at $Q = 5$.

(The sample program output is on the next page, and the Python 3 source code is at the end of this document. →)

Sample program output

[Homework 8 Problems 2–4]

SVM with soft margins ($C=0.01$) and polynomial kernel ($Q=2$):

classifier	number of support vectors	in-sample error	out-of-sample error
0 vs. all	2179	0.105884	0.111609
1 vs. all	386	0.014401	0.021923
2 vs. all	1970	0.100261	0.098655
3 vs. all	1950	0.090248	0.082711
4 vs. all	1856	0.089425	0.099651
5 vs. all	1585	0.076258	0.079721
6 vs. all	1893	0.091071	0.084704
7 vs. all	1704	0.088465	0.073244
8 vs. all	1776	0.074338	0.082711
9 vs. all	1978	0.088328	0.088191

[Homework 8 Problems 5–6]

SVM with soft margins ($C=1$) and polynomial kernel ($Q=5$) for 1 vs. 5 classifier:

C	Q	number of support vectors	in-sample error	out-of-sample error
0.0001	2.0	236.0	0.008969	0.016509
0.0010	2.0	76.0	0.004484	0.016509
0.0100	2.0	34.0	0.004484	0.018868
0.1000	2.0	24.0	0.004484	0.018868
1.0000	2.0	24.0	0.003203	0.018868
0.0001	5.0	26.0	0.004484	0.018868
0.0010	5.0	25.0	0.004484	0.021226
0.0100	5.0	23.0	0.003844	0.021226
0.1000	5.0	25.0	0.003203	0.018868
1.0000	5.0	21.0	0.003203	0.021226

Cross Validation

In the next two problems, we will experiment with 10-fold cross validation for the polynomial kernel. Because E_{cv} is a random variable that depends on the random partition of the data, we will try 100 runs with different partitions and base our answer on how many runs lead to a particular choice.

7. Consider the 1 versus 5 classifier with $Q = 2$. We use E_{cv} to select $C \in \{0.0001, 0.001, 0.01, 0.1, 1\}$. If there is a tie in E_{cv} , select the smaller C . Within the 100 random runs, which of the following statements is correct?

Answer: [b] $C = 0.001$ is selected most often.

8. Again, consider the 1 versus 5 classifier with $Q = 2$. For the overall winning selection in the previous problem, the average value of E_{cv} over the 100 runs is closest to

Answer: [c] 0.005

Sample program output

[Homework 8 Problems 7-8]

Cross-validation error for soft margin ($C=1$) SVM with polynomial kernel ($Q=2$) for 1 vs. 5 classifier:

C	cross-validation error	selection rate
0.0001	0.009711	0.00
0.0010	0.004734	0.52
0.0100	0.004695	0.21
0.1000	0.004772	0.11
1.0000	0.004785	0.16

RBF Kernel

Consider the radial basis function (RBF) kernel $K(\mathbf{x}_n, \mathbf{x}_m) = \exp(-\|\mathbf{x}_n - \mathbf{x}_m\|^2)$ in the soft-margin SVM approach. Focus on the 1 versus 5 classifier.

9. Which of the following values of C results in the lowest E_{in} ?

Answer: [e] $C = 10^6$

10. Which of the following values of C results in the lowest E_{out} ?

Answer: [c] $C = 100$

Sample program output

[Homework 8 Problems 9-10]

Soft margin SVM with RBF kernel for 1 vs. 5 classifier:

C	in-sample error	out-of-sample error
0.01	0.003844	0.023585
1.00	0.004484	0.021226
100.00	0.003203	0.018868
10000.00	0.002562	0.023585
1000000.00	0.000641	0.023585

(The Python 3 source code is on the following page. →)

Python 3 source code

```
from pathlib import Path

import numpy as np
import pandas as pd
import requests
from sklearn import svm
from sklearn.model_selection import cross_val_score
from sklearn.utils import shuffle

DATA_DIR = Path(__file__).parents[2] / "data"

if __name__ == "__main__":
    rng = np.random.default_rng()

    DATA_DIR.mkdir(exist_ok=True)
    data = {}
    for dataset in ["train", "test"]:
        file = f"features.{dataset}"
        if not (DATA_DIR / file).exists():
            r = requests.get(f"http://www.amlbook.com/data/zip/{file}")
            with open(DATA_DIR / file, "wb") as f:
                f.write(r.content)
        data[dataset] = np.loadtxt(DATA_DIR / file)

    C = 0.01
    Q = 2
    clf = svm.SVC(C=C, kernel="poly", degree=Q, gamma=1, coef0=1)
    df = pd.DataFrame(columns=["classifier", "number of support vectors",
                              "in-sample error", "out-of-sample error"])

    for digit in range(10):
        x_train = data["train"][:, 1:]
        y_train = 2 * (data["train"][:, 0] == digit) - 1
        clf.fit(x_train, y_train)
        df.loc[digit] = (
            f"{digit} vs. all",
            clf.n_support_.sum(),
            1 - clf.score(x_train, y_train),
            1 - clf.score(data["test"][:, 1:],
                          2 * (data["test"][:, 0] == digit) - 1)
        )
    print("\n[Homework 8 Problems 2-4]\n"
          f"Soft margin ({C=}) SVM with polynomial kernel ({Q=}): \n",
          df.to_string(index=False), sep="")
```

```

x_train = data["train"][np.isin(data["train"][:, 0], (1, 5))]
y_train = 2 * (x_train[:, 0] == 1) - 1
x_test = data["test"][np.isin(data["test"][:, 0], (1, 5))]
y_test = 2 * (x_test[:, 0] == 1) - 1
df = pd.DataFrame(columns=["C", "Q", "number of support vectors",
                           "in-sample error", "out-of-sample error"])

for Q in (2, 5):
    for C in (Cs := (0.0001, 0.001, 0.01, 0.1, 1)):
        clf = svm.SVC(C=C, kernel="poly", degree=Q, gamma=1, coef0=1)
        clf.fit(x_train[:, 1:], y_train)
        df.loc[len(df)] = (
            C, Q, clf.n_support_.sum(),
            1 - clf.score(x_train[:, 1:], y_train),
            1 - clf.score(x_test[:, 1:], y_test)
        )
print("\n[Homework 8 Problems 5-6]\n"
      f"Soft margin ({C=}) SVM with polynomial kernel ({Q=}) for "
      "1 vs. 5 classifier:\n",
      df.to_string(index=False), sep="")

Q = 2
N_runs = 100
N_folds = 10
clfs = [svm.SVC(C=C, kernel="poly", degree=Q, gamma=1, coef0=1)
         for C in Cs]
counters = np.zeros((2, len(Cs)), dtype=float)
for _ in range(N_runs):
    Es_cv = tuple(1 - cross_val_score(clf, x_train[:, 1:], y_train,
                                       cv=N_folds).mean()
                  for clf in clfs)
    counters[0] += Es_cv
    counters[1, np.argmax(Es_cv)] += 1
    x_train, y_train = shuffle(x_train, y_train)
counters /= N_runs
df = pd.DataFrame({"C": Cs, "cross-validation error": counters[0],
                  "selection rate": counters[1]})
print("\n[Homework 8 Problems 7-8]\n"
      f"Cross-validation error for soft margin ({C=}) SVM with "
      f"polynomial kernel ({Q=}) for 1 vs. 5 classifier:\n",
      df.to_string(index=False), sep="")

df = pd.DataFrame(columns=["C", "in-sample error", "out-of-sample error"])
for C in (0.01, 1, 100, 1e4, 1e6):
    clf = svm.SVC(C=C, gamma=1)
    clf.fit(x_train[:, 1:], y_train)
    df.loc[len(df)] = (
        clf.C,
        1 - clf.score(x_train[:, 1:], y_train),
        1 - clf.score(x_test[:, 1:], y_test)
    )
print("\n[Homework 8 Problems 9-10]\n"
      "Soft margin SVM with RBF kernel for 1 vs. 5 classifier:\n",
      df.to_string(index=False), sep="")

```